

# Assignment 2

Mohammad Saad Khan - 1004138818

March 21, 2020

## 1 Implementing the Model

```
function log_prior(zs)
    return factorized_gaussian_log_density(0, 0, zs)
end

function logp_a_beats_b(za,zb)
    return -log1pexp(zb - za)
end


function all_games_log_likelihood(zs,games)
    zs_a = zs[games[:,1],:]
    zs_b = zs[games[:,2],:]
    likelihoods = logp_a_beats_b.(zs_a, zs_b)
    return sum(likelihoods, dims=1)
end

function joint_log_density(zs,games)
    return all_games_log_likelihood(zs,games) + log_prior(zs)
end
```


## 2 Examining the Posterior

```
plot(title="Prior Gaussian Contour Plot",
      xlabel = "Player 1 Skill",
      ylabel = "Player 2 Skill"
)
prior_gaussian(zs) = exp(log_prior(zs))
skillcontour!(prior_gaussian)
plot_line_equal_skill!()
```

## Prior Gaussian Contour Plot




```
plot(title="Likelihood Gaussian Contour Plot",  
      xlabel = "Player 1 Skill",  
      ylabel = "Player 2 Skill"  
    )  
likelihood_gaussian(zs) = exp(all_games_log_likelihood(zs, two_player_toy_games(1, 1)))  
skillcontour!(likelihood_gaussian)  
plot_line_equal_skill!()
```




```
plot(title="Joint Gaussian Contour Plot - A wins 1",
      xlabel = "Player 1 Skill",
      ylabel = "Player 2 Skill"
    )
joint_gaussian_A1(zs) = exp(joint_log_density(zs, two_player_toy_games(1, 0)))
skillcontour!(joint_gaussian_A1)
plot_line_equal_skill!()
```

## Joint Gaussian Contour Plot - A wins 1




```
plot(title="Joint Gaussian Contour Plot - A wins 10",  
      xlabel = "Player 1 Skill",  
      ylabel = "Player 2 Skill"  
    )  
joint_gaussian_A10(zs) = exp(joint_log_density(zs, two_player_toy_games(10, 0)))  
skillcontour!(joint_gaussian_A10)  
plot_line_equal_skill!()
```

## Joint Gaussian Contour Plot - A wins 10



```
plot(title="Joint Gaussian Contour Plot - Both win 10",  
      xlabel = "Player 1 Skill",  
      ylabel = "Player 2 Skill"  
    )  
joint_gaussian_A10B10(zs) = exp(joint_log_density(zs, two_player_toy_games(10, 10)))  
skillcontour!(joint_gaussian_A10B10)  
plot_line_equal_skill!()
```

## Joint Gaussian Contour Plot - Both win 10



## 3 Stochastic Variational Inference

```
function elbo(params, logp, num_samples)
    N = length(params[1])
    samples = (randn(N, num_samples) .* exp.(params[2])) .+ params[1]
    logp_estimate = logp(samples)
    logq_estimate = factorized_gaussian_log_density(params[1], params[2], samples)
    return mean(logp_estimate - logq_estimate)
end

function neg_toy_elbo(params; games = two_player_toy_games(1,0), num_samples = 100)
    logp(zs) = joint_log_density(zs, games)
    return -elbo(params, logp, num_samples)
end

function fit_toy_variational_dist(init_params, toy_evidence; num_itrs=200, lr= 1e-2,
num_q_samples = 10)
    params_cur = init_params
    for i in 1:num_itrs
        grad_params = gradient((params)->neg_toy_elbo(params,
            games=toy_evidence, num_samples=num_q_samples), params_cur)[1]

        params_cur = params_cur .- (lr .* grad_params)
        @info neg_toy_elbo(params_cur, games=toy_evidence, num_samples=num_q_samples)
        plot();
        p(zs) = exp(joint_log_density(zs, toy_evidence))
        skillcontour!(p, colour=:red) # plot likelihood contours for target posterior
        plot_line_equal_skill!()
        q(zs) = exp(factorized_gaussian_log_density(params_cur[1], params_cur[2], zs))
        # display(skillcontour!(q, colour=:blue)) # plot likelihood contours for variational
        posterior
    end
end
```

```

    return params_cur, neg_toy_elbo(params_cur, games=toy_evidence,
num_samples=num_q_samples)
end

num_players_toy = 2
toy_mu = [-2.,3.] # Initial mu, can initialize randomly!
toy_ls = [0.5,0.] # Initial log_sigma, can initialize randomly!
toy_params_init = (toy_mu, toy_ls)


games = two_player_toy_games(1, 0)
fitted_A1_params, loss = fit_toy_variational_dist(toy_params_init, games)
print(loss)

0.8173925834300657

plot(title="Variational Gaussian Contour Plot - A wins 1",
      xlabel = "Player 1 Skill",
      ylabel = "Player 2 Skill"
    );
p(zs) = exp(joint_log_density(zs,games))
skillcontour!(p,colour=:red) # plot likelihood contours for target posterior
plot_line_equal_skill!()
q(zs) = exp(factorized_gaussian_log_density(fitted_A1_params[1], fitted_A1_params[2],
zs))
skillcontour!(q, colour=:blue)

```

## Variational Gaussian Contour Plot - A wins 1



```

games = two_player_toy_games(10, 0)
fitted_A10_params, loss = fit_toy_variational_dist(toy_params_init, games)
print(loss)


```

2.987927859284215

```

plot(title="Variational Gaussian Contour Plot - A wins 10",
     xlabel = "Player 1 Skill",
     ylabel = "Player 2 Skill"
);
p(zs) = exp(joint_log_density(zs,games))
skillcontour!(p,colour=:red) # plot likelihood contours for target posterior
plot_line_equal_skill!()
q(zs) = exp(factorized_gaussian_log_density(fitted_A10_params[1], fitted_A10_params[2],
zs))
skillcontour!(q, colour=:blue)

```



```

games = two_player_toy_games(10, 10)
fitted_A10B10_params, loss = fit_toy_variational_dist(toy_params_init, games)
print(loss)

```

15.666049643567101


```

plot(title="Variational Gaussian Contour Plot - Both win 10",
     xlabel = "Player 1 Skill",
     ylabel = "Player 2 Skill"
);
p(zs) = exp(joint_log_density(zs,games))
skillcontour!(p,colour=:red) # plot likelihood contours for target posterior
plot_line_equal_skill!()
q(zs) = exp(factorized_gaussian_log_density(fitted_A10B10_params[1],
fitted_A10B10_params[2], zs))
skillcontour!(q, colour=:blue)

```



## Variational Gaussian Contour Plot - Both win 10



## 4 Approximate Inference

(a) Yes, the games of other players besides  $i$  and  $j$  provides information about their skills

```
function fit_variational_dist(init_params, tennis_games; num_itrs=200, lr= 1e-2,
num_q_samples = 10)
    params_cur = init_params
    for i in 1:num_itrs
        grad_params = gradient((params)->neg_toy_elbo(params,
            games=tennis_games, num_samples=num_q_samples), params_cur)[1]

        params_cur = params_cur .- (lr .* grad_params)
        @info neg_toy_elbo(params_cur, games=tennis_games, num_samples=num_q_samples)
    end
    return params_cur, loss
end
```

Fitting distribution to observed dataset


```
init_mu = randn(length(player_names))
init_log_sigma = randn(length(player_names))
init_params = (init_mu, init_log_sigma)

# Train variational distribution
trained_params, loss = fit_variational_dist(init_params, tennis_games)
print(loss)
```

15.666049643567101

Plotting approximate mean and variance of all players

```
perm = sortperm(trained_params[1])
plot(trained_params[1][perm], yerror=exp.(trained_params[2][perm]))
```



Printing top 10 players by mean skill


```
print(player_names[perm][end-9:end])
```

```
Any["Richard-Gasquet", "Juan-Martin-Del-Potro", "Tomas-Berdych", "Jo-Wilfri  
ed-Tsonga", "Robin-Soderling", "David-Ferrer", "Andy-Murray", "Rafael-Nadal  
", "Roger-Federer", "Novak-Djokovic"]
```

Plotting joint distribution over Nadal and Federer

```
rog_index = findall(x->x=="Roger-Federer", player_names)
raf_index = findall(x->x=="Rafael-Nadal", player_names)

plot(legend=:topleft,
      title="Approximate Posterior of Skills of Federer and Nadal",
      xlabel = "Roger Federer Skill",
      ylabel = "Rafael Nadal Skill"
    );
means = [trained_params[1][rog_index]; trained_params[1][raf_index]]
logsigs = [trained_params[2][rog_index]; trained_params[2][raf_index]]
dist(zs) = exp.(factorized_gaussian_log_density(means, logsigs, zs))
skillcontour!(dist)
plot_line_equal_skill!()
```



To find exact probability that one player has higher skill than another. Consider two players with skills  $z_A$  and  $z_B$  and from a distribution with mean  $\mu$  and variance  $\Sigma$ . Firstly, consider the matrix  $A = \begin{bmatrix} 1 & -1 \\ 0 & 1 \end{bmatrix}$ .

As in the hint, we can use this matrix transformation to create new variables  $y = Az$ . We know that  $y_A = z_A - z_B$  and  $y_B = z_B$  because of the matrix above.

Also, because  $A$  is a linear transformation and  $z \sim N(\mu, \Sigma)$ , we know  $y \sim N(A\mu, A\Sigma A^T)$  because  $y = Az$ .

Then, we specifically want to find  $P(z_A - z_B > 0) = P(y_A > 0)$ . Since we are using a factorized gaussian distribution, we know  $y_A \sim N(y_A, [A\Sigma A^T]_{11})$

We can use the reparameterization trick to express this as  $P(y_A + [A\Sigma A^T]_{11}N > 0)$  where  $N$  is the standard normal.

Finally,  $P(y_A + [A\Sigma A^T]_{11}N > 0) = P[A\Sigma A^T]_{11}N > -y_A) = 1 - P(N \leq \frac{-y_A}{[A\Sigma A^T]_{11}})$

```
function transform_skills(mu, logsig)
    A = [1 -1; 0 1]
    ya = mu[1] - mu[2]
    sig = Diagonal(exp.(logsig))
    cov = A * sig * A'
    var = sqrt(cov[1])
    threshold = -ya/var
    return 1 - cdf(Normal(0, 1), threshold)
end

function monte_carlo(mu, logsigs)
    num_samples = 10000
    samples = (randn(2, num_samples) .* exp.(logsigs)) .+ mu
    diff = samples[1,:] - samples[2,:]
end
```

```

    return count(x -> x > 0, diff)/num_samples
end

```

First, we can find the probability that Federer is better than Nadal:

```
print(transform_skills(means, logsigs))
```

```
0.5416807784390549
```

```
print(monte_carlo(means, logsigs))
```

```
0.5818
```

Then, we can compare Federer to the worst player

```

means = [trained_params[1][rog_index]; trained_params[1][perm[1]]]
logsigs = [trained_params[2][rog_index]; trained_params[2][perm[1]]]

```

```
print(transform_skills(means, logsigs))
```

```
0.9999911763364223
```

```
print(monte_carlo(means, logsigs))
```

```
1.0
```

If we changed the prior to  $\text{Normal}(10, 1)$ , then the sections that would change are (b), (c), (e), (g), (h)